


Batch-one GPU latency



Peak CUDA memory

