

Accuracy vs practical CPU inference cost

